PLAIN-ENGLISH GUIDE

DeepSeek Harness

What it is, where you can use it, and how you can shape it for your own applications.

Think of it as an engine room for an AI assistant

The Harness gives an AI agent the pieces it needs to work: a model, memory, tools, safety checks, approvals, and ways to interact with people.

Its main idea is simple: **almost everything is replaceable**. You can add, remove, or swap parts without rebuilding the whole assistant.

Where we are today

Ready for learning and internal experiments

Installed locally, built successfully, checked, documented, and supported by a searchable brain.

Still a developer-preview project

Expect changes. It should not yet be treated as a finished, stable product for critical production work.

How it works

You ask
A goal or task

›

The agent plans
Uses memory and rules

›

Tools act
With checks and approval

The Harness keeps a reviewable history, prepares information for the model, controls available tools, and records the results.

USEFUL APPLICATIONS

Where can you use it?

The best fit is an AI assistant that needs tools, memory, control, or several interchangeable parts.

1

Coding assistant

Read files, suggest changes, run checks, keep a task history, and ask before sensitive actions.

2

Research and review assistant

Collect evidence, compare sources, produce reports, and preserve what supports each conclusion.

3

Internal operations assistant

Guide repeatable workflows, prepare actions, check rules, and leave important decisions to a person.

4

Tool-using chat assistant

Connect a conversational assistant to files, search, business services, or your own application functions.

5

Multi-agent work

Give bounded tasks to specialist agents, collect their results, and keep one owner responsible for the final answer.

6

AI feature inside an existing app

Reuse selected ideas, such as tool approvals or reliable history, without moving your entire application into the Harness.

Best use for you right now

Use the current Harness as a reference laboratory. Study how it solves agent memory, tools, approvals, and modular design. Then bring only the useful pieces into each existing application.

When it is probably too much

If you only need a simple chatbot or one model call, a smaller direct integration will usually be easier.

CUSTOMIZATION

How can you modify it?

Treat it like a set of building blocks. Change one layer at a time and keep the safety boundaries visible.

YOU WANT TO CHANGE	WHAT YOU WOULD CUSTOMIZE
The AI model	Swap or add a model provider while keeping the rest of the assistant unchanged.
What the agent can do	Add tools for your files, APIs, databases, products, or internal workflows.
Safety and control	Choose which actions need permission, which files are available, and what should be refused.
Memory	Decide what is recorded, how conversations are reconstructed, and where history is stored.
The experience	Use a terminal, Web interface, headless service, or connect the agent to your own app.
Specialist behavior	Create focused plugins or subagents for coding, research, marketing, operations, or review.

Best parts to reuse first

1. **Reliable history:** keep a clear record of what the user, model, and tools did.
2. **One controlled tool path:** every action passes through the same checks.
3. **Human approval:** ask at the moment a sensitive action would happen.
4. **Replaceable services:** change providers without changing every feature that uses them.
5. **Clean shutdown:** stop background work and remove temporary capabilities safely.

Important limits

The current local version is **demo-verified**, not market-ready. Browser use, live model-provider behavior, other operating systems, long-running production use, and a full security review still need separate testing.

A sensible next step

Pick one existing application. Identify one problem, such as unreliable agent history or unsafe tool execution. Port that small piece first, test it inside the real application, and only then consider the next layer.